CHANDU PRASAD BHAIRAPU



AI Engineer

bhairapuchanduprasad@gmail.com **(+49) 15560194877** **Frankfurt • EU Blue Card** **linkedin** **github**

ABOUT ME

Master's graduate in Artificial Intelligence with 3 years of professional experience in AI engineering and software development. Specialized in full-stack AI applications, from multi-agent LLM pipelines and RAG to React and Python services deployed with Docker and CI/CD on AWS and Azure. Most recently sole developer of an agentic knowledge-extraction system for the German energy sector, working side by side with domain experts and the CTO.

EXPERIENCE

Associate Generative AI Developer - Agentic Indexing Pipeline Engineer for German Energy Sector

Kniit GmbH

November 2025 – present Muldestausee, Germany

- Sole developer owning an agentic AI knowledge-extraction product end to end, from data ingestion, chunking and prompt engineering to indexing (agentic, deterministic and monolithic strategies), working directly with the CTO and SAP domain experts across the German energy sector.
- Developed Expert Interface, a graph curation web app (React frontend, Python microservices backend) deployed on Azure, enabling domain experts to validate and approve knowledge-graph changes.
- Designed and implemented a Model Context Protocol (MCP) query service exposing the knowledge graph for millisecond-latency retrieval, integrated directly into Claude Code for agentic development workflows.
- Designed and optimized multi-agent extraction algorithms with LangGraph (entity-extraction, relationship-analysis, consolidation and reflection agents running in parallel) for complex regulatory and technical documents, raising precision/recall from 33%/45% to 93%/90% against an SAP-system-derived ground truth.
- Engineered a deterministic (code-based) extraction pipeline for a second, high-volume source stream, cutting a $200,000 projected LLM budget to near-zero recurring cost while pushing extraction accuracy toward 100%.
- Built software as an AI-augmented engineer with coding agents (Claude Code, custom skills, agentic coding workflows), applying clean code principles.
- **Tech stack: Python, React, FastAPI, LangChain, LangGraph, LiteLLM, MCP, RAG, GraphRAG, Vector Search, Postgres, MongoDB, Neo4j, Docker, Kubernetes, Azure**

GenAI Developer (AI Consulting Intern) - Development and Deployment of an Agentic AI Workflow

Covestro Deutschland AG

December 2024 – April 2025 Leverkusen, Germany

- Designed and deployed an AI agent with LangGraph that routes each request across vector database queries, general LLM inference and summarization, after mapping the decision-making workflows of business stakeholders.
- Developed a React front-end and a Python backend with FastAPI, using Postman for API testing.
- Built and containerized the entire stack with Docker and deployed it to AWS using CDK, following clean code practices; version-controlled in GitLab with CI/CD.
- **Tech stack: Python, LangGraph, LangChain, FastAPI, React, Postman, Docker, AWS (CDK, EC2, Lambda), GitLab CI/CD**

GenAI Developer (Master's Thesis) - Evaluating the Performance of RAG Systems

Covestro Deutschland AG

May 2024 – November 2024 Leverkusen, Germany

- Designed Retrieval-Augmented Generation (RAG) pipelines, focusing on chunking strategies, vector databases, indexing, and fine-tuning models for optimized retrieval.
- Automated RAG generation and retrieval evaluation with RAGAS and PytrecEval, comparing model answers to chemist-provided ground truth.
- Benchmarked multilingual embedding models for vector search on over 1M data chunks of Covestro's internal chemical data.
- Built AWS cloud infrastructure with Lambda, SQS and DynamoDB for PrismaChat RAG evaluation using CDK.

- **Tech stack: Python, LangChain, Hugging Face Transformers, PyTorch, RAGAS, PytrecEval, Vector Databases, Azure AI Search, AWS (Lambda, SQS, DynamoDB, CDK), Docker, Streamlit**

Data Scientist (Intern) – ML-Based Prediction of Prepolymer Synthesis Outcomes

Covestro Deutschland AG

November 2023 – April 2024 | Leverkusen, Germany

- Built and compared Random Forest, XGBoost and LSTM models to predict prepolymer synthesis outcomes (e.g. viscosity, reaction progression) from time-series process data, evaluating tree-based vs. deep learning approaches for this tabular/temporal hybrid problem.
- Engineered lag features and rolling statistical windows from raw sensor time-series (temperature, reaction time, process variables) to feed the tree-based models, while feeding raw sequences directly into the LSTM.
- Benchmarked model performance against traditional empirical curve-fitting methods, improving early identification of batches likely to deviate from target specification.
- Deployed an interactive Streamlit dashboard exposing model predictions to chemists in real time, letting them simulate parameter changes and see predicted outcomes before running physical batches.
- **Tech stack: Python, Scikit-Learn, XGBoost, TensorFlow/Keras, Pandas, NumPy, SciPy, Matplotlib, Streamlit**

Machine Learning Engineer - Electrical Vehicle Load Forecasting using Time Series Analysis

Cognizant Technology Solutions

September 2021 – March 2022 | Hyderabad, India

- Developed ARIMA, SARIMA, XG-Boost regressor, FB Prophet, LSTM models for time series analysis for EVs.
- Developed LSTM, achieved 96% test accuracy and strong performance on unseen data.
- **Tech stack: Python, Scikit-Learn, TensorFlow**

PROJECTS

Project: Ops Assistant | TypeScript | Node.js | React | PostgreSQL | MCP | GitHub Actions
- Internal assistant on synthetic operations data: a TypeScript (Node.js) MCP server over PostgreSQL, a React + TypeScript UI, unit and end-to-end tests, and an automated evaluation gate in GitHub Actions.

Project: Agent Eval & Tracing Harness | Python | Langfuse | Docker | GitHub Actions
- Evaluation harness for an LLM agent calling a mocked business API: Langfuse tracing, a golden-question set and automated regression checks on every pull request.

TECHNICAL SKILLSET

- **Full-stack Applications & Backend:** Python, FastAPI, React, Python microservices, REST APIs, Postman, SQL, Postgres, NoSQL (MongoDB, Neo4j), vector databases (Qdrant, Pinecone, Azure AI Search)
- **LLM & Agentic AI:** LLM frameworks (LangChain, LiteLLM), agent frameworks (LangGraph), MCP, RAG, GraphRAG, prompt engineering, evaluation frameworks (RAGAS, PytrecEval), Hugging Face Transformers, PyTorch
- **Cloud & DevOps:** Docker, Kubernetes, Azure, AWS (S3, EC2, Lambda, SQS, DynamoDB, CDK), Git, GitLab CI/CD, coding agents (Claude Code), Agile
- **Languages:** English (Business Fluent), German (B1), Hindi (Fluent), Telugu (Native)

EDUCATION

Master of Science in Artificial Intelligence — March 2022 – April 2025
Technische Hochschule Würzburg-Schweinfurt, Würzburg, Germany

Bachelor of Technology in Mechanical Engineering (Minor in Computer Science & Engineering) — August 2017 – July 2021
G Pulla Reddy Engineering College, Kurnool, India

PUBLICATIONS & CERTIFICATIONS

- **[1]** Navigating the Future: An approach of Autonomous Indoor Vehicles
- Udemy - LLMs Mastery : Transformers & Generative AI